

```
In [ ]: import numpy as np
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import train_test_split

import nltk
from nltk import word_tokenize
from nltk.stem import WordNetLemmatizer, PorterStemmer
from nltk.corpus import wordnet
```

```
In [ ]: nltk.download("wordnet")
nltk.download('punkt')
nltk.download('averaged_perceptron_tagger')
nltk.download('omw-1.4')
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Unzipping corpora/wordnet.zip.
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data]   /root/nltk_data...
[nltk_data]   Unzipping taggers/averaged_perceptron_tagger.zip.
```

```
Out[ ]: True
```

```
In [ ]: # https://www.kaggle.com/shivamkushwaha/bbc-full-text-document-classification
!wget -nc https://lazyprogrammer.me/course_files/nlp/bbc_text_cls.csv
```

```
--2021-11-24 05:54:58-- https://lazyprogrammer.me/course_files/nlp/bbc_text_cls.csv
Resolving lazyprogrammer.me (lazyprogrammer.me)... 104.21.23.210, 172.67.213.166, 260
6:4700:3030::ac43:d5a6, ...
Connecting to lazyprogrammer.me (lazyprogrammer.me)|104.21.23.210|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 5085081 (4.8M) [text/csv]
Saving to: 'bbc_text_cls.csv'
```

```
bbc_text_cls.csv 100%[=====>] 4.85M 19.9MB/s in 0.2s
```

```
2021-11-24 05:54:59 (19.9 MB/s) - 'bbc_text_cls.csv' saved [5085081/5085081]
```

```
In [ ]: df = pd.read_csv('bbc_text_cls.csv')
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	text	labels
0	Ad sales boost Time Warner profit\n\nQuarterly...	business
1	Dollar gains on Greenspan speech\n\nThe dollar...	business
2	Yukos unit buyer faces loan claim\n\nThe owner...	business
3	High fuel prices hit BA's profits\n\nBritish A...	business


	text	labels
--	------	--------

```
In [ ]: len(df)
```

```
Out[ ]: 2225
```

```
In [ ]: inputs = df['text']  
labels = df['labels']
```

```
In [ ]: labels.hist(figsize=(10, 5));
```



```
In [ ]: inputs_train, inputs_test, Ytrain, Ytest = train_test_split(  
    inputs, labels, random_state=123)
```

```
In [ ]: vectorizer = CountVectorizer()
```

```
In [ ]: Xtrain = vectorizer.fit_transform(inputs_train)  
Xtest = vectorizer.transform(inputs_test)
```

```
In [ ]: Xtrain
```

```
Out[ ]: <1668x26287 sparse matrix of type '<class 'numpy.int64'>'  
    with 337411 stored elements in Compressed Sparse Row format>
```

```
In [ ]: (Xtrain != 0).sum()
```

```
Out[ ]: 337411
```

```
In [ ]: # what percentage of values are non-zero?
        (Xtrain != 0).sum() / np.prod(Xtrain.shape)
```

```
Out[ ]: 0.007695239935415004
```

```
In [ ]: model = MultinomialNB()
        model.fit(Xtrain, Ytrain)
        print("train score:", model.score(Xtrain, Ytrain))
        print("test score:", model.score(Xtest, Ytest))
```

```
train score: 0.9922062350119905
test score: 0.9712746858168761
```

```
In [ ]: # with stopwords
        vectorizer = CountVectorizer(stop_words='english')
        Xtrain = vectorizer.fit_transform(inputs_train)
        Xtest = vectorizer.transform(inputs_test)
        model = MultinomialNB()
        model.fit(Xtrain, Ytrain)
        print("train score:", model.score(Xtrain, Ytrain))
        print("test score:", model.score(Xtest, Ytest))
```

```
train score: 0.9928057553956835
test score: 0.976606822262118
```

```
In [ ]: def get_wordnet_pos(treebank_tag):
        if treebank_tag.startswith('J'):
            return wordnet.ADJ
        elif treebank_tag.startswith('V'):
            return wordnet.VERB
        elif treebank_tag.startswith('N'):
            return wordnet.NOUN
        elif treebank_tag.startswith('R'):
            return wordnet.ADV
        else:
            return wordnet.NOUN
```

```
In [ ]: class LemmaTokenizer:
        def __init__(self):
            self.wnl = WordNetLemmatizer()
        def __call__(self, doc):
            tokens = word_tokenize(doc)
            words_and_tags = nltk.pos_tag(tokens)
            return [self.wnl.lemmatize(word, pos=get_wordnet_pos(tag)) \
                    for word, tag in words_and_tags]
```

```
In [ ]: # with lemmatization
vectorizer = CountVectorizer(tokenizer=LemmaTokenizer())
Xtrain = vectorizer.fit_transform(inputs_train)
Xtest = vectorizer.transform(inputs_test)
model = MultinomialNB()
model.fit(Xtrain, Ytrain)
print("train score:", model.score(Xtrain, Ytrain))
print("test score:", model.score(Xtest, Ytest))
```

train score: 0.9922062350119905
test score: 0.9676840215439856

```
In [ ]: class StemTokenizer:
    def __init__(self):
        self.porter = PorterStemmer()
    def __call__(self, doc):
        tokens = word_tokenize(doc)
        return [self.porter.stem(t) for t in tokens]
```

```
In [ ]: # with stemming
vectorizer = CountVectorizer(tokenizer=StemTokenizer())
Xtrain = vectorizer.fit_transform(inputs_train)
Xtest = vectorizer.transform(inputs_test)
model = MultinomialNB()
model.fit(Xtrain, Ytrain)
print("train score:", model.score(Xtrain, Ytrain))
print("test score:", model.score(Xtest, Ytest))
```

train score: 0.9892086330935251
test score: 0.9694793536804309

```
In [ ]: def simple_tokenizer(s):
    return s.split()
```

```
In [ ]: # string split tokenizer
vectorizer = CountVectorizer(tokenizer=simple_tokenizer)
Xtrain = vectorizer.fit_transform(inputs_train)
Xtest = vectorizer.transform(inputs_test)
model = MultinomialNB()
model.fit(Xtrain, Ytrain)
print("train score:", model.score(Xtrain, Ytrain))
print("test score:", model.score(Xtest, Ytest))
```

train score: 0.9952038369304557
test score: 0.9712746858168761

```
In [ ]: # What is the vector dimensionality in each case?
# Compare them and consider why they are larger / smaller
```